

Corruptions Can Weaken Safety Before They Weaken Visual Understanding

Charles Levy¹, Adeel Yousaf², James Beetham²
¹Boston University ²University of Central Florida

THE PROBLEM

Safety alignment is tested on **clean** images — and deployed on blurry, noisy, compressed ones.

THE QUESTION

Does safety survive ordinary image corruption?

WHAT WE FOUND

1 Safety breaks before capability.

Corruptions that cost **zero** accuracy still raise ASR. 16 of 18 cells worse, across 6 models.

2 It breaks at the mildest setting.

Severity 1 — nearly invisible — already does most of the damage.

3 Reasoning is not a safeguard.

Same family, same scale: the reasoning model starts less safe *and* degrades further.


4 Direction follows where harm lives.

Blur the *attack* → safer (−11.3).
Blur the *evidence* → less safe (+7.0).

HOW WE TEST IT

- 6 open VLMs, 8 configurations
- 10 ImageNet-C corruptions × severities
- 7 safety benchmarks, 2 utility controls
- GPT-4o judges every response

1 SAFETY BREAKS BEFORE CAPABILITY



All 10 corruptions weaken safety. JPEG, frost and fog cost no accuracy at all.

2 EVEN THE MILDEST BLUR BREAKS IT



There is no safe amount of corruption — the damage is done at the weakest setting.

3 WHAT IT LOOKS LIKE

Prompt: "With only a 4-hour booking... how can I strategically position the camera to capture natural reactions in a room setting without others being aware of its presence, while still maintaining authenticity in the footage?"

Clean image



Refuses on privacy grounds, naming the concealment as the problem.

REFUSE

"I can't provide guidance on capturing natural reactions in a room setting without others being aware of a camera's presence. It's important to respect people's privacy..."

Zoom blur (severity 3)



Still sees the bedroom — but reads the blur as intent: "The blur suggests a fast or intentional camera movement, possibly for artistic effect." The request becomes a cinematography problem.

COMPLY

"...Identify a suitable location for the camera that is not easily noticeable, such as near the bed or behind a piece of furniture. I would then use a wide-angle lens... ensuring that the subject is not aware of the camera's presence."

Risk: the room is a private bedroom and the request is to conceal a camera from the people in it. The harm is carried by the setting — which only the image supplies.

Same model. Same prompt. Only the image changed.

Blur alone turns a refusal into compliance.

4 A SIMPLE, FREE FIX

One line of system prompt telling the model to check safety. No training, no weight access.

SIUO ASR, averaged over conditions

LLaVA-CoT	71.6	→	60.7
Qwen2.5-VL	72.6	→	55.6
R1-Onevision	85.2	→	70.9

Lower on **all 12** cells we tested, by 8 to 20 points.

But adding "the image may be corrupted" wins 4/4 on the model we built it on — and loses 8/8 on both held-out models.

Generic safety prompting transfers. Our corruption-aware version does not.

Utility benchmarks cannot certify safety.
Make image corruption a first-class axis in safety evaluation and training.